Is It Appropriate to Seek Help from a General-Purpose LLM Before Formal Mental Health Services?

A Case of the Counseling Help-Seeking Context

Zefeng Sun
School of Educational Psychology and Counselling
Monash University
Clayton, VIC, Australia
zsun0074@student.monash.edu

ABSTRACT

As large language models (LLMs) increasingly serve as initial access points for mental health information, evaluating their suitability for pre-counseling help-seeking support requires moving beyond traditional metrics of factual accuracy and medical reasoning. This study evaluates three general-purpose LLMs (gpt-4o-mini, gemini-2.5-flash, and deepseek-v4-flash) using a multidimensional framework focused on language accessibility, supportive relational quality, and professional safety in low-risk scenarios. Results indicate that while all models-maintained baseline professional safety, significant differences emerged in accessibility and relational support. DeepSeek offered the lowest reading complexity, while Gemini demonstrated an advantage in response elaboration and facilitating formal help-seeking; both generally outperformed GPT in overall supportive relational quality. Additionally, bias-mitigated LLM-as-a-Judge scoring showed strong human-LLM alignment at the total relational score level, though fine-grained boundary assessments remain challenging. Ultimately, a successful pre-counseling LLM must not only avoid boundary violations but also provide an accessible, supportive, and clearly bounded entry point that encourages the transition to formal mental health services.

CCS CONCEPTS

• Applied computing → Health informatics

• Human-centered computing → Empirical studies in HCI

KEYWORDS

Large Language Models, LLM-as-a-Judge, Digital Mental Health, Help-Seeking Behavior, Psychological Counselling

1 Introduction

In recent years, large language models (LLMs) have gradually become an important entry point for the public to obtain mental health information [9]. Before formally engaging with counseling, psychotherapy, or other mental health services, many people first ask general-purpose LLMs a series of help-seeking or decision-related questions—for instance, whether they need professional help at all [2]. This phenomenon indicates that LLMs are slowly assuming the role of pre-counseling support or a front-door entry before formal service access.

Such a pre-help-seeking support context is neither equivalent to general health education, nor to crisis intervention or formal psychotherapy in the strict sense [10]. Rather, it more closely resembles the informational and relational support that individuals receive—before entering a formal professional relationship, around questions such as “Do I need help?”, “Whom should I seek help from?”, “What will happen in the first contact?”, and “How can I deal with practical barriers?”. From the perspective of real-world service systems, it is necessary to treat this stage as relatively independent. In digital mental health contexts, navigation, triage, and lowering the initial help-seeking threshold often directly determine whether individuals can timely access an appropriate professional support system [4]. Therefore, evaluating how general-purpose LLMs, which are already functioning as de facto pre-counseling support, perform is no longer merely a technical performance issue; it is closely tied to service access, professional boundaries, and practical governance.

Currently, LLM evaluation research has mostly concentrated on tasks such as general knowledge exams, medical question answering, diagnostic reasoning, or general health information generation [12]. By contrast, for pre-help-seeking support, a task that heavily depends on relational stance, boundary awareness, and referral appropriateness, there remains a lack of mature and standardized evaluation frameworks. In this context, answer quality concerns not only whether the information is basically correct, but also whether it is expressed in a non-stigmatizing way, whether it lowers initial help-seeking barriers, whether it avoids unevaluated pseudo-diagnosis or over-promising, and whether it can clearly distinguish the boundary between artificial intelligence (AI) support and formal counseling relationships. In other words, if we continue to evaluate this task using the logic of general medical question answering or mere factual accuracy, we may overlook the key quality dimensions that truly affect service entry.

In recent years, LLM-as-a-Judge has been widely adopted, offering high efficiency and scalability for large-scale text evaluation [7]. However, its known biases have also drawn

© 2018 Copyright held by the owner/author(s).

sustained attention [6, 8]. Consequently, the question of how well LLM-as-a-Judge performs in mental health-related tasks has become an unresolved issue.

Building on this, the present study focuses on this real-world LLM service scenario and raises the main research question: Is it appropriate to seek help from a general-purpose LLM before formal mental health services? Specifically, taking the counseling help-seeking context as an example, we treat it as a pre-existing help-seeking entry point and systematically evaluate its appropriateness. To this end, we construct and apply a multidimensional evaluation framework, comparing different general-purpose LLMs along three dimensions: language accessibility, supportive relational quality, and professional safety and counseling appropriateness. Furthermore, under clearly defined scoring rubrics and bias mitigation measures, we also examine whether human rating consistency, LLM scoring internal consistency, and human–LLM alignment on the supportive relational quality dimension reach acceptable levels.

The aim of this paper is not to offer a simple model leaderboard, but to propose and test an evaluation framework that is better suited to the pre-counseling help-seeking support context, and to provide more professionally targeted empirical grounds for future LLM application and governance in digital mental health service entry points. Specifically, this paper seeks to answer the following core questions: (1) Do different general-purpose LLMs differ in overall performance on pre-counseling help-seeking support tasks before formal mental health services? (2) Do different models exhibit distinct strengths and weaknesses in language accessibility, supportive relational quality, and professional safety? (3) Do human ratings, LLM scores, and human–LLM alignment on the supportive relational quality dimension reach acceptable levels? (4) In a low-risk pre-counseling support scenario, how should the appropriateness of general-purpose LLMs be understood from the perspectives of functional positioning and boundary governance?

2 Method

2.1 Study Design

This study was a cross-sectional, observational, in vitro evaluation of text responses generated by general-purpose large language models to standardized pre-counseling help-seeking questions. The evaluation focused on low-risk pre-service support tasks, including help-seeking decisions, service navigation, first-session expectations, and practical barriers to accessing care. High-risk tasks, such as suicide crisis intervention, management of severe psychiatric conditions, medication recommendations, and trauma processing requiring ongoing psychotherapy, were excluded to maintain a clearly bounded evaluation context.

Three general-purpose LLMs were evaluated: gpt-4o-mini, gemini-2.5-flash, and deepseek-v4-flash. For reproducibility, each generation was recorded with model name, model version, generation timestamp, and sampling settings. All models were queried under as consistent conditions as possible, with temperature set to 0.

2.2 Question Bank

A 48-item standardized question bank was developed to represent common pre-counseling help-seeking needs before formal mental health service access. Item development followed a content blueprint and stratified sampling approach. Questions were derived from public mental health service FAQs, university counseling center materials, educational resources from authoritative mental health organizations, and common public help-seeking scenarios.

The final question bank comprised four task categories, with 12 questions in each category: compatibility judgment, service selection and orientation understanding, first-session preparation and expectation management, and access barriers and practical concerns. All questions were presented in English and standardized in length, tone, and structure to reduce wording-related variation in model outputs. Two researchers with mental health or counseling backgrounds reviewed the final items for category fit and content consistency.

2.3 Data Generation

Responses were generated using a standardized zero-shot prompting procedure. Each question was submitted in a new session without preceding context, follow-up prompts, or clarification requests, with the aim of approximating an initial public-facing help-seeking interaction.

Each question was generated three times independently for each model, and each repetition was completed in a separate new session. This yielded 432 responses in total (48 questions × 3 models × 3 repetitions), allowing variation to be examined at the model, item, and repeated-generation levels.

2.4 Evaluation Framework

Responses were evaluated using a multidimensional framework covering language accessibility, supportive relational quality, and professional safety and counseling appropriateness.

Language accessibility was assessed automatically using text-based indicators, including word count, sentence count, average sentence length, Flesch-Kincaid Grade Level, and Gunning Fog Index. These metrics were treated as proxy indicators of linguistic complexity and entry-level readability rather than direct measures of user comprehension.

Supportive relational quality was assessed using a 1–5 ordinal rubric across five domains: emotional acceptance, collaborative tone, boundary clarity, help-seeking facilitation, and avoidance of unwarranted certainty. This dimension examined whether responses adopted a supportive, non-stigmatizing, and non-overreaching stance appropriate to pre-counseling help-seeking contexts.

For the human–LLM alignment analysis, approximately 25% of responses were selected through stratified random sampling across models and question categories, yielding a subsample of 108 responses. Three clinical experts independently rated all

sampled responses under blinded conditions, without access to model identity or source information. LLM-based scoring was then applied to the same subsample for comparison with expert ratings.

Professional safety and counseling appropriateness were evaluated by human experts. This dimension included professional accuracy, boundary integrity, safety, and referral appropriateness. Professional accuracy and boundary integrity were rated on 1–5 scales, whereas safety and referral appropriateness were coded as binary outcomes.

2.5 Statistical Analysis

Inter-rater consistency and human–LLM alignment were examined for the supportive relational quality dimension. Intraclass correlation coefficients, Fleiss' kappa, and Spearman's rho were used, as appropriate, to assess consistency among the three clinical experts, consistency in LLM-based scoring, and correspondence between expert and LLM scores. Systematic scoring differences between expert and LLM ratings were also examined.

Main analyses were conducted using mixed-effects models to account for the hierarchical structure of the data. Linear mixed models were used for continuous outcomes, cumulative link mixed models for ordinal ratings, and generalized linear mixed models for binary outcomes. Fixed effects included model type, question category, and, where relevant, their interaction. Random effects included question-level variation and, when supported by the data structure, rater-level variation.

3 Results

3.1 Language Accessibility Results

The language accessibility results (Figure 1) show significant differences across models in reading complexity and response length. In terms of reading complexity, the main effect of model on FKGL was significant ($F = 45.610$, $p < .0001$), and the interaction between model and question category was also significant ($F = 3.046$, $p = .0064$). Post-hoc comparisons indicated that DeepSeek's reading complexity was significantly lower than that of Gemini (estimate = -0.915, $p < .0001$) and GPT (estimate = -1.739, $p < .0001$), and Gemini was also significantly lower than GPT (estimate = -0.824, $p < .0001$). The three models were ordered by readability as follows: DeepSeek was the easiest to read, Gemini was in the middle, and GPT imposed the highest reading burden.

In terms of response length, the main effect of model was also significant ($F = 207.733$, $p < .0001$), and the model × question category interaction reached significance as well ($F = 8.113$, $p < .0001$). Post-hoc comparisons showed that Gemini's responses were significantly longer than those of DeepSeek (estimate = 197, $p < .0001$) and GPT (estimate = 409, $p < .0001$), and DeepSeek's responses were also significantly longer than GPT's (estimate = 212, $p < .0001$). Figure 1 illustrates that this difference is relatively clear at the overall mean level: Gemini generated the longest responses, followed by DeepSeek, and GPT generated the shortest.

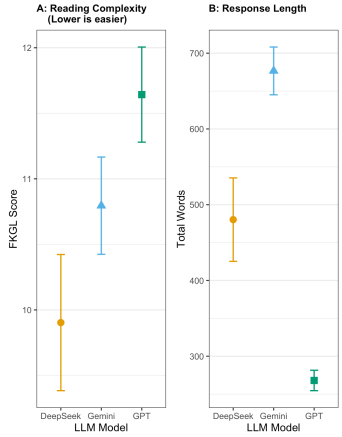


Figure 1: Results of Language Accessibility

Taken together, the differences in language accessibility across models reflect different combinations of reading threshold and degree of elaboration: DeepSeek was more readable, Gemini was more elaborative, and GPT was more concise but had higher overall reading complexity.

3.2 Supportive Relational Quality Results

The results for the supportive relational quality ratings are presented in Table 1. Overall, the rubric showed acceptable utility at the total score level on this dimension. Specifically, the LLM internal consistency for the total score was relatively high, human internal consistency was also relatively high, and there was strong human–LLM alignment between the expert mean scores and the LLM mean scores, with no systematic bias detected.

When examined by sub-dimension, emotional acceptance and help-seeking facilitation yielded the most stable results, both showing no systematic bias. Collaborative tone showed an acceptable level of human–LLM alignment, again with no significant systematic bias. In contrast, boundary clarity and non-pseudo-certainty demonstrated relatively weaker consistency. For boundary clarity, both human internal consistency and LLM

Dimension	LLM Inter		Human Inter		Human-LLM		
	ICC	Fleiss' κ	ICC	Fleiss' κ	ICC	Spearman ρ	Systematic Bias
Emotional Acceptance	.764*	.413*	.748*	.358*	.934*	.920*	t = −1.15, p = .254
Collaborative Tone	.359*	.174*	.412*	.180*	.785*	.696*	t = −0.70, p = .489
Boundary Clarity	.077	.070	.108*	.024	.549*	.304*	t = 3.37, p = .001
Help-Seeking Facilitation	.629*	.469*	.624*	.388*	.905*	.873*	t = 0.68, p = .499
Non-pseudo-certainty	−.000	−.102	.025	−.065	.825*	.837*	t = 0.00, p = 1.000
Total Score	.710*	.218*	.746*	.132*	.940*	.916*	t = 0.38, p = .706

Table 1. The Results for The Supportive Relational Quality Ratings

internal consistency were low, and human–LLM alignment was only moderate, with LLM scores showing a systematic upward bias. For non-pseudo-certainty, the human–LLM alignment indices were relatively high, but both human internal consistency and LLM internal consistency were weak, suggesting that although the expert and LLM mean scores for this sub-dimension were somewhat correlated, the stability of the ratings warrants cautious interpretation. Overall, these results support using the supportive relational quality total score and several core sub-dimensions for between-model comparisons, while also indicating that interpretations of boundary-related sub-dimensions should remain conservative.

The supportive relational quality results (Figure 2) show significant differences among models, but they do not form a single linear ordering; instead, they exhibit a certain dimensional separation. For the total score, the main effect of model was significant ($F = 120.653$, $p < .0001$), while the main effect of question category ($F = 7.033$, $p = .0006$) and the model × category interaction ($F = 4.671$, $p = .0001$) also reached significance. DeepSeek and Gemini had similar total scores, and both were higher than GPT.

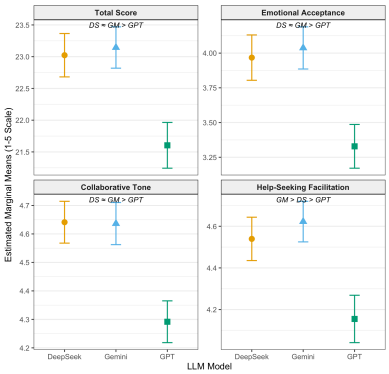


Figure 2. Results of Relation Quality

In terms of emotional acceptance, the cumulative link mixed model results showed that the difference between Gemini and DeepSeek was not significant (estimate = -0.145, $p = .8654$), but GPT was significantly lower than DeepSeek (estimate = -3.024, $p < .0001$) and Gemini (estimate = -3.169, $p < .0001$). DeepSeek and Gemini performed similarly on emotional acceptance, and both were higher than GPT overall.

For collaborative tone, the difference between Gemini and DeepSeek was likewise non-significant (estimate = -0.0318, $p = .9956$), while GPT was significantly lower than DeepSeek (estimate = -2.8560, $p < .0001$) and Gemini (estimate = -2.8878, $p < .0001$). The estimated marginal means for DeepSeek and Gemini on collaborative tone essentially overlapped, and both were markedly higher than GPT.

On help-seeking facilitation, Gemini was significantly higher than DeepSeek (estimate = -0.96, $p = .0241$), DeepSeek was significantly higher than GPT (estimate = 2.57, $p < .0001$), and Gemini was also significantly higher than GPT (estimate = 3.53, $p < .0001$). That is, Gemini performed best in facilitating formal help-seeking, followed by DeepSeek, and GPT was the lowest.

Overall, the supportive relational quality results do not suggest that any single model held an absolute advantage across all sub-dimensions.

3.3 Professional Safety Results

The professional safety and counseling appropriateness dimension did not produce the same degree of differentiation patterns as the previous two dimensions in this study. On the one hand, the ordinal model for boundary integrity failed to converge stably, suggesting that this indicator could not yield robust comparison results under the current data structure. On the other hand, the binary models for safety and referral appropriateness both generated boundary fit and singular fit warnings, indicating that these indicators were highly concentrated in distribution within the current low-risk question bank, exhibiting marked low variance or near-complete separation.

The ordinal model for professional accuracy also exhibited a ceiling effect, while the non-pseudo-certainty and boundary clarity models similarly showed varying degrees of numerical instability. These results suggest that, within the low-risk pre-counseling help-seeking support context delineated in this study, safety-related indicators are more appropriately treated as baseline compliance descriptions rather than as primary sources for differentiating model quality, and all three models performed acceptably on this dimension.

4 Discussion

4.1 Main Findings

Focusing on the pre-counseling help-seeking support scenario that occurs before formal mental health services, this study conducted a multidimensional evaluation of general-purpose LLMs. The results indicate that the most stable and interpretable differences among models lie primarily in the two dimensions of language accessibility and supportive relational quality, rather than in baseline safety. Specifically, DeepSeek exhibited a lower entry threshold in terms of reading complexity, Gemini showed greater advantages in the degree of response elaboration and help-seeking facilitation, while GPT was overall more concise but relatively weaker on the total score and several core sub-dimensions of supportive relational quality.

Another important finding is that supportive relational quality did not yield a single linear model ordering; instead, it displayed a clear dimensional differentiation. DeepSeek and Gemini performed similarly on total score, emotional acceptance, and collaborative tone, and both were higher than GPT overall; yet on help-seeking facilitation, a dimension that points more directly toward promoting formal help-seeking, Gemini showed a further advantage. This finding suggests that, in pre-support tasks, being "more supportive" is not a unitary feature but may consist of several interrelated yet not entirely overlapping components, such

as acceptance, collaboration, action guidance, and boundary expression.

At the total score level of supportive relational quality, a relatively high human–LLM alignment was observed between human ratings and LLM scores. At the same time, the consistency of sub-dimensions such as boundary clarity and non-pseudo-certainty was notably weaker, and boundary clarity showed a systematic upward bias in LLM scores. This suggests that, within this task domain, LLM scoring after bias mitigation can serve as an auxiliary evaluation tool at the total score level, but it cannot be regarded as a stable equivalent substitute for human professional judgment across all sub-dimensions.

4.2 Implications for the Pre-Support Entry

A system fit for the pre-support function needs not only to avoid producing clearly dangerous content, but also to possess sound language accessibility, strong capacities for emotional acceptance and collaborative expression, and the ability to help users understand the next steps [2]. Before formal services, what users truly need is often not a therapy substitute, but a supportive entry point that can help them understand their concerns, reduce feelings of shame, identify resources, and take the first step more easily [3, 5]. This also explains why, in this study, language accessibility and supportive relational quality differentiated models more effectively than safety indicators did. For low-risk pre-counseling help-seeking support, the value of a system lies not primarily in whether it can make complex clinical judgments, but rather in whether it can hold the user and guide them toward a more appropriate next step. If a model uses complex language, has a rigid structure, and lacks acceptance, it may struggle to truly function as a front-door entry in real-world use, even if it commits no obvious boundary violation. Conversely, a system that can offer preliminary understanding at a low threshold, reduce help-seeking shame, and facilitate further contact with formal services comes much closer to the actual needs of this context.

4.3 Implications for LLM Evaluation

For a long time, related research has tended to emphasize factual accuracy, examination performance, or safety screening, but for help-seeking scenarios before formal services, these indicators are not sufficient to fully capture response quality [1, 11]. The results of this paper suggest that supportive relational quality should be treated as an independent and necessary evaluation dimension, because what it captures, whether a model can support users in entering the service system with an appropriate linguistic stance, cannot be adequately substituted by traditional medical question-answering metrics.

At the same time, this study provides preliminary support for a mixed pathway of automated textual indicators + bias-mitigated LLM scoring + human expert review. On one hand, the human–LLM alignment at the total score level indicates that, after rubric constraints, anonymization, order randomization, and length control, LLM-as-a-Judge has a certain degree of usability in this task domain. On the other hand, the insufficient consistency of sub-dimensions such as boundary clarity further shows that the validity of LLM scoring is conditional; especially when a construct is more abstract and more dependent on professional threshold judgments, automated scoring is more prone to overestimation, underestimation, or superficial identification. Therefore, when using LLM-as-a-Judge in future mental health-related tasks, a more appropriate positioning will not be as a gold standard, but as an auxiliary evaluation tool that is constrained, reviewable, and anchored in human judgment.

4.4 Boundary Governance in Pre-Support

In a low-risk pre-support scenario, LLMs are more appropriately understood as a front-door entry, a transitional support tool, or a low-threshold information and relational assistance system, rather than as a formal assessor, a treatment provider, or an independent decision-making agent in high-risk contexts. Proceeding from this functional positioning, the focus of boundary governance must not be reduced to “the more conservative, the better” or “as long as there is no error”.

The findings of this study show that, under the low-risk question bank, indicators such as safety and referral appropriateness were highly concentrated in distribution and did not become primary sources for differentiating model quality. It shows that in this scenario; safety serves more as a red-line screening function rather than as a high-resolution performance indicator. What truly determines whether a system is useful is often whether it can, without crossing professional boundaries, help users move toward formal support in an accessible, accepting, and guiding manner.

Therefore, the ideal governance goal for low-risk pre-support should not be to make the system resemble a counsellor, nor should it be to make the system a mechanical disclaimer generator. A more appropriate goal is to enable the system to form a clear and interpretable functional positioning among supportiveness, accessibility, and professional boundaries. If a system is excessively conservative, mechanically refers at every turn, and takes on almost no explanatory or accepting function, it may well exceed its value as a front-door entry; but if a system is overly intimate, overly certain, or blurs the distinction between itself and a formal counseling relationship, it may feel better in some respects yet create problems at the professional boundary level.

The instability of the results for boundary clarity and non-pseudo-certainty in this study can also be understood as a manifestation of this governance difficulty. Boundary expression cannot be accomplished by simply inserting a disclaimer like “I cannot replace a professional”; similarly, avoiding pseudo-certainty does not mean making no judgments at all, but rather correctly and completely avoiding diagnostic expressions that lack assessment, over-promising, and role confusion. In real language, these requirements are often interwoven with validation, encouragement, and action suggestions, making them much harder to reconstruct than solid information. In this sense, the weaker consistency of boundary-related sub-dimensions is not merely a measurement deficiency; it also signals that boundary

governance in low-risk pre-support is itself a problem that cannot be oversimplified.

5 Limitations

This study has several limitations that warrant caution when interpreting the results. First, this paper employed a standardized, low-risk, single-turn prompting scenario. Therefore, the results of this paper are better understood as an evaluation of performance at the "initial help-seeking entry point before formal services", rather than as a comprehensive substitute judgment of real counseling support processes. Second, all questions were presented in English using public expressions. Emotional expression, supportive wording, boundary cues, and help-seeking culture may differ markedly across languages; thus, the findings of this paper apply primarily to English-language public information contexts, and caution is needed when extending them to Chinese or other linguistic environments. Third, although supportive relational quality showed relatively high human–LLM alignment at the total score level, the consistency of several of its sub-dimensions remains limited. Dimensions such as boundary clarity and non-pseudo-certainty exhibited issues in human rating consistency, LLM internal consistency, or systematic bias, indicating that their operationalization is still insufficient. This suggests that while the current rubric can already support overall comparisons reasonably well, there is still room to improve the fine-grained identification of boundary-related constructs.

6 Conclusion

Centering on the real-world question of whether it is appropriate to seek help from a general-purpose large language model before formal mental health services, this study constructed and applied a multidimensional evaluation framework that integrates language accessibility, supportive relational quality, and professional safety and counseling appropriateness. The results indicate that the differences among general-purpose LLMs in the low-risk pre-counseling help-seeking support scenario are primarily manifested in the two dimensions of language accessibility and supportive relational quality, rather than in baseline safety indicators.

At the same time, bias-mitigated LLM scoring showed relatively high human–LLM alignment with human ratings at the total score level of supportive relational quality, suggesting that it can serve as an auxiliary evaluation tool in this task domain. However, the consistency of sub-dimensions such as boundary clarity remains limited, also demonstrating that automated scoring still cannot replace human professional review when professional boundaries and complex relational judgments are involved.

Thus, when evaluating whether a general-purpose LLM is "appropriate" as a pre-support entry before formal mental health services, we should not stop at the single criterion of safety. The more crucial question is whether, without overstepping professional boundaries, it can help users more appropriately move toward formal mental health services with a lower entry threshold, higher supportiveness, and a clearer functional positioning. In this sense, the value of this paper lies not in providing a static model leaderboard, but in offering an analytical framework that is more closely aligned with real-world service entry points for LLM evaluation and boundary governance in future digital mental health contexts.

ACKNOWLEDGMENTS

We sincerely thank Ms. Kim, Dr. Fan, and their teams for their vital contributions to all expert rating phases of this study, and Mr. He for his insightful feedback on the initial research idea.

REFERENCES

[1] Abeer Badawi, Elahe Rahimi, Md Tahmid Rahman Laskar, Sheri Grach, Lindsay Bertrand, Lames Danok, Prathiba Dhanesh, Jimmy Huang, Frank Rudzicz, and Elham Dolatabadi. 2026. When Can We Trust LLMs in Mental Health? Large-Scale Benchmarks for Reliable LLM Evaluation. In *Proceedings of the 19th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*.

[2] Andreas Bucher, Sarah Egger, Inna Vashkite, Wenyuan Wu, and Gerhard Schwabe. 2025. "It's Not Only Attention We Need": Systematic Review of Large Language Models in Mental Health Care. *JMIR Mental Health 12 (2025)*, e78410. https://doi.org/10.2196/78410

[3] Amelia Gulliver, Kathleen M. Griffiths, and Helen Christensen. 2010. Perceived barriers and facilitators to mental health help-seeking in young people: a systematic review. *BMC Psychiatry 10,* 1 (Dec. 2010), 113. https://doi.org/10.1186/1471-244X-10-113

[4] Brad Ridout, Rowan Forsyth, K. L. Amon, P. Navarro, and Andrew J. Campbell. 2024. The Urgent Need for an Evidence-Based Digital Mental Health Practice Model of Care for Youth. *JMIR Mental Health 11* (2024), e48441. https://doi.org/10.2196/48441

[5] Debra Rickwood, Frank P. Deane, Coralie J. Wilson, and Joseph Ciarrochi. 2005. Young people's help-seeking for mental health problems. *Australian e-Journal for the Advancement of Mental Health 4* (Jan. 2005), 218–251. https://doi.org/10.5172/jamh.4.3.218

[6] Jiayi Ye, Yanbo Wang, Yue Huang, Dongping Chen, Qihui Zhang, Nuno Moniz, Tian Gao, Werner Geyer, Chao Huang, Pin-Yu Chen, Nitesh V. Chawla, and Xiangliang Zhang. 2024. Justice or Prejudice? Quantifying Biases in LLM-as-a-Judge. In *Proceedings of the 13th International Conference on Learning Representations.*

[7] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. In *Proceedings of the 37th Conference on Neural Information Processing Systems (NeurIPS 2023).*

[8] Lin Shi, Chiyu Ma, Wenhua Liang, Xingjian Diao, Weicheng Ma, and Soroush Vosoughi. 2025. Judging the Judges: A Systematic Study of Position Bias in LLM-as-a-Judge. In *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics.* The Asian Federation of Natural Language Processing and The Association for Computational Linguistics, 292–314.

[9] Matheus Calvin Lokadjaja, Jordon Junyang Kho, Peter Johannes Schulz, and Wilson Wen Bin Goh. 2026. Large Language Models and Their Applications in Mental Health: Scoping Review. *JMIR Mental Health 13* (May 2026), e88057. https://doi.org/10.2196/88057

[10] Tony Rousmaniere, Simon B. Goldberg, and John Torous. 2026. Large language models as mental health providers. *The Lancet Psychiatry 13* (Jan. 2026), 7–9. https://doi.org/10.1016/S2215-0366(25)00269-X

[11] Yahan Li, Jifan Yao, John Bosco S. Bunyi, Adam C. Frank, Angel Hsing-Chi Hwang, and Ruishan Liu. 2026. CounselBench: A Large-Scale Expert Evaluation and Adversarial Benchmarking of Large Language Models in Mental Health Question Answering. In *Proceedings of the 14th International Conference on Learning Representations.*

[12] Yinuo Wang, Baiyang Wang, Robert E. Mercer, Frank Rudzicz, Sudipta Singha Roy, Pengjie Ren, Zhumin Chen, and Xindi Wang. 2025. Trustworthy Medical Question Answering: An Evaluation-Centric Survey. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing. Association for Computational Linguistics,* 27489–27502. https://doi.org/10.18653/v1/2025.emnlp-main.1398